

CS2043 - Unix Tools & Scripting

Lecture 4

More Unix Tools

Spring 2015 ¹

Instructor: Nicolas Savva

January 28, 2015

¹based on slides by Hussam Abu-Libdeh, Bruno Abrahao and David Slater over the years

Course Logistics

- Last day to enroll!
- Assignment 1 is due tonight
- Late policy: total of 5 days for the course
- Additional OH and support (beginning next week)

- Accessing Remote Resources
- Variables
- More useful commands
- Piping, input/output redirection

ssh

You can use “secure shell” (ssh) to connect to a remote machine.

```
ssh [username@]<remote machine name or IP address>
```

- If the username is omitted, local username will be used.
- Remote machine has to be configured to accept ssh connections:
 - ssh daemon (service) has to be running and listening on an open port (by default 22)

Executing remote commands

ssh can be used to execute commands on the remote machine

Example

```
ssh nsavva@csug01.csuglab.cornell.edu ls
```

this will execute `ls` on `csug01.csuglab.cornell.edu` and output the result to the screen before `ssh` terminates the connection

- You can use the `-f` flag to put `ssh` into the background before executing the remote command
- You can use the `-Y` flag to forward X11 (graphical windows/user interface) to the local machine

Run firefox on the remote machine

```
ssh -Y nsavva@csug01.csuglab.cornell.edu firefox
```

Identity files

- You can use an identity file to authenticate with the remote machine instead of using your username/password.
- An identity file allows you to authenticate yourself with a “*pass phrase*” (which could be empty).
- Identity files are typically a pair of public/private keys used for asymmetric key cryptography (e.g., RSA).

Identity files

To use identity files,

- 1 Create an identity file using `ssh-keygen`

create identity files using RSA encryption

```
ssh-keygen -t rsa
```

- 2 Append the generated public key file (by default `~/.ssh/id_rsa.pub`) to the `~/.ssh/authorized_keys` file on the remote machine.
- 3 `ssh` to the remote machine using the `-i` flag to use the identity file, and specify the private file corresponding to the public file appended at the remote machine

```
ssh nsavva@example.com -i ~/.ssh/id_rsa
```

ssh configuration file

You can configure ssh to use customized settings when connecting to a particular host without having to set the corresponding flags every time. The file `~/.ssh/config` contains these settings.

Sample config

```
host rgblab
hostname maxwell.cs.cornell.edu

host tesla
hostname tesla.cs.cornell.edu
user nsavva
ForwardX11 yes
IdentityFile ~/.ssh/id_rsa
```

Here, `ssh rgblab` connects to `maxwell.cs.cornell.edu` and `ssh tesla` connects to `tesla.cs.cornell.edu` with username `nsavva` and identity `~/.ssh/id_rsa` and enable X11 forwarding.

sftp

- Transfer files securely between local and remote machines.
- Operates over an encrypted ssh transport.
 - same connection settings as ssh
- Uses an interactive console to interact with the user
 - unless the `-b [batchfile]` option is used to use batch files
- Useful sftp commands:
 - `help` : to see a list of commands and help on them
 - `put` : upload a file to the remote machine
 - `get` : download a file from the remote machine
 - `cd` / `pwd` : change directory / print current directory on remote machine
 - `lcd` / `lpwd` : change directory / print current directory on local machine

scp

- Copy files securely over a network using an encrypted ssh transport.

copy file to remote machine

```
scp file nsavva@remote_machine:
```

copy file from remote machine

```
scp nsavva@remote_machine:file .
```

- The ':' is necessary after the remote machine name. A path on the remote machine starting from the user's home directory can be specified after the colon ':'.

copy directories using the -r flag

```
scp -r pics_dir nsavva@remote_machine:
```

Other useful commands

wget

```
wget [OPTIONS] [URL...]
```

Download a file from a remote location over HTTP. Popular options:

- `-r` : recursive
- `-c` : continue a partial download

curl

```
curl [OPTIONS] [URL...]
```

Transfer data from/to web servers.

For more info on these commands, consult the `man` pages.

Variables

- Bash scripting is powerful (you could write a web server just using bash scripting)
- We need variables to really get anything done
- All variables preceded by the dollar sign (\$)
- The contents of any variable can be listed using the echo command
- Two types of variables: Environment and Local

Example:

```
echo My shell is $SHELL and the username is $USER  
My shell is /bin/zsh and the username is nsavva
```

Environment Variables

- Environment Variables are used by the system to define aspects of operation
- The shell passes environment variables to its child processes
- Examples:
 - \$SHELL : which shell will be used by default
 - \$PATH : a list of directories to search for binaries
 - \$HOSTNAME : the hostname of the machine
 - \$HOME : the home directory for the current user
- To get a list of all current environment variables use the `env` command

New Environment Variable:

Set a new environment variable using `export`

```
nsavva@x200t:~$ export X=42
```

```
nsavva@x200t:~$ echo $X
```

```
42
```

Local Variables

We can define local variables which only exist in the current shell:

New Environment Variable:

Set a new environment variable using `export`

```
nsavva@x200t:~$ x=7
```

```
nsavva@x200t:~$ echo $x
```

```
7
```

Note: You cannot have a space after the `x` nor before the `7`

- The main difference between environment and local variables is that the environment variables are passed to child processes while local variables are not.
- A copy is passed (variable changes in the child processes are not reflected in parent)
- We will talk more about this in a few lectures

Listing and Removing Variables

- `env` : displays all environment variables
- `set` : displays all shell/local variables
- `unset name` : remove a shell variable
- `unsetenv name` : remove an environment variable

WC

- How many lines of code are in my new awesome program?
- How many words are in this document?
- Good for bragging rights

Word, Character, Line, and Byte count with `wc`

- `wc -l` : count the number of lines
- `wc -w` : count the number of words
- `wc -m` : count the number of characters
- `wc -c` : count the number of bytes

sort

Sorts the lines of a text file alphabetically.

- `sort -r -u file`
 - sorts the file in reverse order and deletes duplicate lines.
- `sort -n -k 2 -t : file`
 - sorts the file numerically by using the second column, separated by a colon

Example

Consider a file (`numbers.txt`) with the numbers 1, 5, 8, 11, 62 each on a separate line, then:

```
$ sort numbers.txt
```

```
1
```

```
11
```

```
5
```

```
62
```

```
8
```

```
$ sort numbers.txt -n
```

```
1
```

```
5
```

```
8
```

```
11
```

```
62
```

uniq

- `uniq file` - Discards all but one of successive identical lines
- `uniq -c file` - Prints the number of successive identical lines next to each line

The Translate Command

```
tr [options] <set1> [set2]
```

- Translate or delete characters
- Sets are strings of characters
- By default, searches for strings matching set1 and replaces them with set2

Example:

```
cat somefile | tr 'AEIOU' 'aeiou' - changes all capital  
vowels to lower case vowels
```

Some Simple Examples

Example:

`echo *` prints everything in the directory, separated by spaces.

Let's separate them by newlines instead:

- `echo * | tr ' ' '\n'` – replaces all spaces with newlines

Example:

Let's print a file in all uppercase:

- `tr 'a-z' 'A-Z' < test.txt` - prints the contents of `test.txt` in all caps

Pipes and redirection

- `tr` only receives input from *standard input* (`stdin`)
 - i.e. keyboard input
- What if we want to operate on files?
 - 1 Piping: `cat somefile | tr 'AEIOU' 'aeiou'`
 - 2 Input redirection: `tr 'AEIOU' 'aeiou' < somefile`

Pipes and input/output redirection are important and useful throughout UNIX.


Applications in UNIX are associated with Input/Output (I/O) Streams:

- #0 : Standard input stream; STDIN
(usually keyboard)
- #1 : Standard output stream; STDOUT
(usually terminal console)
- #2 : Standard error stream; STDERR
(depends on system setting, but usually terminal console)

UNIX Philosophy

In UNIX you will find many tools that specialize in one or a few things, and they do them really well! To get more complex functionality combine one or more tools by piping or I/O redirection

Standard Streams




Bash scripting is all about combining simple commands together to do more powerful things. This is accomplished using the "pipe" character

Piping

```
<command1> | <command2>
```

- Passes the output from command1 to input of command2
- Works for lots of programs that take input and provide output to the terminal

Piping Streams



Piping Example

Example:

```
ls -al /bin | less
```

- Allows you to scroll through the long list of programs in /bin

```
history | tail -20 | head -10
```

- Displays the 10th-19th last commands from the current session

Redirection

To redirect Input/Output streams, use one of `>` `>>` `<`
Input/Output Streams

- to redirect standard input, use the `<` operator
`command < file`
- to redirect standard output, use the `>` operator
`command > file`
- to redirect standard error, use the `>` operator and specify the stream by number (2)
`command 2> file`

Combining streams

You can combine two streams together by using `2>&1`

This says: send standard error to where standard output is going.
Useful for debugging/catching error messages.

Redirection example

Bash processes I/O redirection from left to right, allowing us to do fun things like this:

Example:

Let's delete everything but the numbers from test1.txt, then store them in test2.txt

- `tr -cd '0-9' < test1.txt > test2.txt`

Starting a Job in the background

To run a job in the background, we will use a new command-line operator:

&

`<command> [arguments] &`

- Runs the specified command as a background job
- Unless told otherwise, will send output to the terminal!

Since `cat` runs indefinitely with no arguments, this will illustrate our point:

Example:

`cat &`

- Try it without the `&`!

Dealing with Excess Output

Many programs output continuously as they run. For example `ping` and `play` both clutter up the terminal with output even when they are backgrounded.

- The solution is to use output redirection

Example:

```
ping google.com > testping.log &
```

- When you care about a program's output, redirect it to a log file.

Example:

```
play somesong.mp3 > /dev/null &
```


- If the text output doesn't matter, redirect it to `/dev/null`.

/dev/null - the black hole

/dev/null is a special file which has the following properties:

- Any user can write to it.
- Anything written to it goes nowhere
- it always reports a successful write.

It works like a black hole for data - you can output to it all day and it will never fill up. Anything you redirect to /dev/null just disappears.



Tee

tee

Redirect your output to a file and still see it on stdout terminal

Example

```
ls -l ~ / | tee homels.txt
```


- Processes and Jobs
- Multiplexing terminals: `tmux` / `screen`
- `find` `grep`, and pattern matching